

Installing MySQL on Mac

1. Download MySQL from Homebrew and Install mysql.

```
rashmi-mbp:~ rashmi$ brew install mysql
--> Downloading https://downloads.sf.net/project/homebrew/bottles/mysql-5.6.19.mavericks.bottle.tar.gz
Already downloaded: /Library/Caches/homebrew/mysql-5.6.19.mavericks.bottle.tar.gz
--> Pouring mysql-5.6.19.mavericks.bottle.tar.gz
--> Caveats
A "/etc/my.cnf" from another install may interfere with a Homebrew-built
server starting up correctly.

To connect:
mysql -uroot

To have launchd start mysql at login:
ln -sfv /usr/local/opt/mysql/*.plist ~/Library/LaunchAgents
Then to load mysql now:
launchctl load ~/Library/LaunchAgents/homebrew.mxcl.mysql.plist
Or, if you don't want/need launchctl, you can just run:
mysql.server start
```

2. Start MySQL server.

```
rashmi-mbp:~ rashmi$ /usr/local/Cellar/mysql/5.6.19/bin/mysql.server start
Starting MySQL
. SUCCESS!
```

Creating Database, Table and loading the data

```
Bye
rashmi-mbp:~ rashmi$ mysql -u root
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 4
Server version: 5.6.19 Homebrew

Copyright (c) 2000, 2014, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> create database nyse;
Query OK, 1 row affected (0.00 sec)

mysql> use nyse;
Database changed
mysql> show tables;
+-----+
| Tables_in_nyse |
+-----+
| nasdaq_daily_prices |
+-----+
1 row in set (0.00 sec)

mysql> desc nasdaq_daily_prices;
+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+
| exchange | varchar(35) | YES | | NULL | |
| stock_symbol | varchar(35) | YES | | NULL | |
| date | varchar(35) | YES | | NULL | |
| stock_price_open | double | YES | | NULL | |
+-----+
```

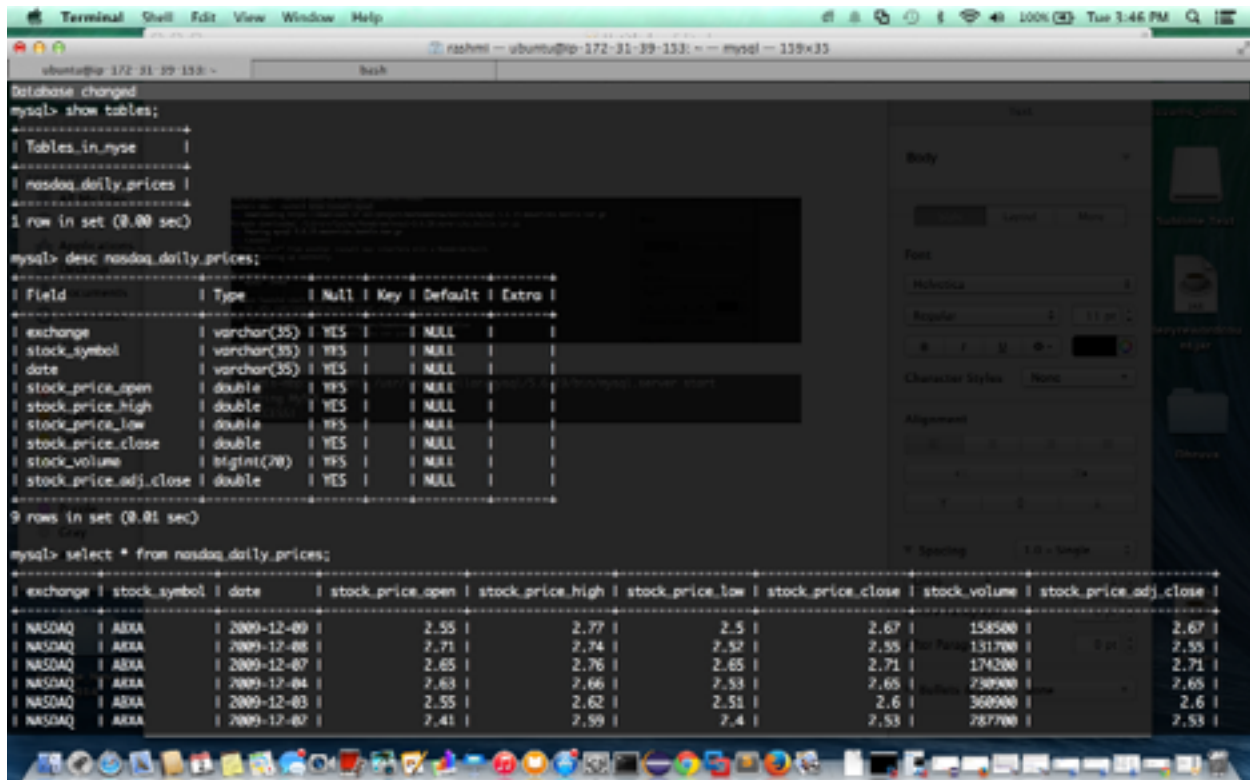
1. Open Mysql command line client, give username as “root” & password. MySQL command line will open. Now we can create tables, databases etc.

2. Create a new database and use new database using command given below.
create database nyse;
use nyse;

3. Create table and insert records

/usr/local/Cellar/mysql/5.6.19/bin/mysql -uroot nyse < tables.sql

/usr/local/Cellar/mysql/5.6.19/bin/mysql -uroot nyse < data.sql



```
Database changed
mysql> show tables;
+-----+
| Tables_in_nyse |
+-----+
| nasdaq_daily_prices |
+-----+
1 row in set (0.00 sec)

mysql> desc nasdaq_daily_prices;
+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+
| exchange | varchar(35) | YES | | NULL | |
| stock_symbol | varchar(35) | YES | | NULL | |
| date | varchar(35) | YES | | NULL | |
| stock_price_open | double | YES | | NULL | |
| stock_price_high | double | YES | | NULL | |
| stock_price_low | double | YES | | NULL | |
| stock_price_close | double | YES | | NULL | |
| stock_volume | bigint(20) | YES | | NULL | |
| stock_price_adj_close | double | YES | | NULL | |
+-----+
9 rows in set (0.01 sec)

mysql> select * from nasdaq_daily_prices;
+-----+
| exchange | stock_symbol | date | stock_price_open | stock_price_high | stock_price_low | stock_price_close | stock_volume | stock_price_adj_close |
+-----+
| NASDAQ | AEXA | 2009-12-09 | 2.55 | 2.77 | 2.5 | 2.67 | 158500 | 2.67 |
| NASDAQ | AEXA | 2009-12-08 | 2.71 | 2.74 | 2.52 | 2.55 | 131700 | 2.55 |
| NASDAQ | AEXA | 2009-12-07 | 2.65 | 2.76 | 2.65 | 2.71 | 174200 | 2.71 |
| NASDAQ | AEXA | 2009-12-04 | 2.63 | 2.66 | 2.53 | 2.65 | 230900 | 2.65 |
| NASDAQ | AEXA | 2009-12-03 | 2.55 | 2.62 | 2.51 | 2.6 | 360900 | 2.6 |
| NASDAQ | AEXA | 2009-12-02 | 2.41 | 2.59 | 2.4 | 2.53 | 287700 | 2.53 |
+-----+
```

Installing Sqoop on Mac

1. Download & untar sqoop-1.4.5.bin_hadoop-0.23.tar.gz

Rashmi-MacBook-Pro:Downloads rashmi\$ tar -zxvf sqoop-1.4.5.bin_hadoop-0.23.tar.gz

2. Setup HADOOP_HOME

```
Rashmis-MacBook-Pro:sqoop-1.4.5.bin__hadoop-0.23 rashmi$ export HADOOP_HOME=/usr/local/Cellar/hadoop/2.6.0/
```

3. Download MySQL connector.

```
Rashmis-MacBook-Pro:lib rashmi$ cp ~/Downloads/mysql-connector-java-5.1.17-bin.jar .
```

Importing data from MySQL to HDFS using Sqoop

1. Execute Scoop command to get the data from MySQL into HDFS.

```
Rashmis-MacBook-Pro:sqoop-1.4.5.bin__hadoop-2.0.4-alpha rashmi$ ./bin/sqoop import --connect jdbc:mysql://localhost/nyse --table nasdaq_daily_prices --target-dir /user/rashmi/nasdaq_daily_prices --username root --split-by exchange
Warning: /Users/rashmi/sqoop-1.4.5.bin__hadoop-2.0.4-alpha/bin/../../hbase does not exist! HBase imports will fail.
Please set $HBASE_HOME to the root of your HBase installation.
Warning: /Users/rashmi/sqoop-1.4.5.bin__hadoop-2.0.4-alpha/bin/../../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /Users/rashmi/sqoop-1.4.5.bin__hadoop-2.0.4-alpha/bin/../../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /Users/rashmi/sqoop-1.4.5.bin__hadoop-2.0.4-alpha/bin/../../zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
14/12/05 21:15:23 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5
14/12/05 21:15:23 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
14/12/05 21:15:23 INFO tool.CodeGenTool: Beginning code generation
14/12/05 21:15:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `nasdaq_daily_prices` AS t LIMIT 1
```

14/12/05 21:15:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.*
FROM `nasdaq_daily_prices` AS t LIMIT 1 [REDACTED]
14/12/05 21:15:23 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/
local/Cellar/hadoop/2.6.0 [REDACTED]
Note: /tmp/sqoop-rashmi/compile/ecd5bff9cfba766fd893f9ff88a4af15/
nasdaq_daily_prices.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details. [REDACTED]
14/12/05 21:15:24 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-rashmi/
compile/ecd5bff9cfba766fd893f9ff88a4af15/nasdaq_daily_prices.jar [REDACTED]
14/12/05 21:15:24 WARN manager.MySQLManager: It looks like you are importing from
mysql. [REDACTED]
14/12/05 21:15:24 WARN manager.MySQLManager: This transfer can be faster! Use the --
direct [REDACTED]
14/12/05 21:15:24 WARN manager.MySQLManager: option to exercise a MySQL-specific
fast path. [REDACTED]
14/12/05 21:15:24 INFO manager.MySQLManager: Setting zero DATETIME behavior to
convertToNull (mysql) [REDACTED]
14/12/05 21:15:24 INFO mapreduce.ImportJobBase: Beginning import of
nasdaq_daily_prices [REDACTED]
14/12/05 21:15:25 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable [REDACTED]
14/12/05 21:15:25 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar [REDACTED]
14/12/05 21:15:25 INFO Configuration.deprecation: mapred.map.tasks is deprecated.
Instead, use mapreduce.job.maps [REDACTED]
14/12/05 21:15:25 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
14/12/05 21:15:27 INFO db.DBInputFormat: Using read committed transaction isolation
14/12/05 21:15:27 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT
MIN(`exchange`), MAX(`exchange`) FROM `nasdaq_daily_prices` [REDACTED]
14/12/05 21:15:27 WARN db.TextSplitter: Generating splits for a textual index column.
14/12/05 21:15:27 WARN db.TextSplitter: If your database sorts in a case-insensitive order,
this may result in a partial import or duplicate records. [REDACTED]
14/12/05 21:15:27 WARN db.TextSplitter: You are strongly encouraged to choose an
integral split column. [REDACTED]
14/12/05 21:15:27 WARN db.BigDecimalSplitter: Set BigDecimal splitSize to
MIN_INCREMENT [REDACTED]
14/12/05 21:15:27 INFO mapreduce.JobSubmitter: number of splits:1 [REDACTED]
14/12/05 21:15:27 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1417842624965_0001 [REDACTED]
14/12/05 21:15:28 INFO impl.YarnClientImpl: Submitted application
application_1417842624965_0001 [REDACTED]
14/12/05 21:15:28 INFO mapreduce.Job: The url to track the job: http://Rashmis-
MacBook-Pro.local:8088/proxy/application_1417842624965_0001/ [REDACTED]

14/12/05 21:15:28 INFO mapreduce.Job: Running job: job_1417842624965_0001
14/12/05 21:15:37 INFO mapreduce.Job: Job job_1417842624965_0001 running in uber mode : false
14/12/05 21:15:37 INFO mapreduce.Job: map 0% reduce 0%
14/12/05 21:15:43 INFO mapreduce.Job: map 100% reduce 0%
14/12/05 21:15:44 INFO mapreduce.Job: Job job_1417842624965_0001 completed successfully
14/12/05 21:15:44 INFO mapreduce.Job: Counters: 30

File System Counters

- FILE: Number of bytes read=0
- FILE: Number of bytes written=114886
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=125
- HDFS: Number of bytes written=538464
- HDFS: Number of read operations=4
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=2

Job Counters

- Launched map tasks=1
- Other local map tasks=1
- Total time spent by all maps in occupied slots (ms)=4634
- Total time spent by all reduces in occupied slots (ms)=0
- Total time spent by all map tasks (ms)=4634
- Total vcore-seconds taken by all map tasks=4634
- Total megabyte-seconds taken by all map tasks=4745216

Map-Reduce Framework

- Map input records=9999
- Map output records=9999
- Input split bytes=125
- Spilled Records=0
- Failed Shuffles=0
- Merged Map outputs=0
- GC time elapsed (ms)=51
- CPU time spent (ms)=0
- Physical memory (bytes) snapshot=0
- Virtual memory (bytes) snapshot=0
- Total committed heap usage (bytes)=149946368

File Input Format Counters

- Bytes Read=0

File Output Format Counters

- Bytes Written=538464

```
14/12/05 21:15:44 INFO mapreduce.ImportJobBase: Transferred 525.8438 KB in 18.9459
seconds (27.755 KB/sec)
14/12/05 21:15:44 INFO mapreduce.ImportJobBase: Retrieved 9999 records.
Rashmis-MacBook-Pro:sqoop-1.4.5.bin__hadoop-2
```

2. Verify the data in HDFS.

```
Rashmis-MacBook-Pro:sqoop-1.4.5.bin__hadoop-2.0.4-alpha rashmi$ hadoop fs -ls /user/
rashmi
14/12/05 21:16:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x - rashmi supergroup      0 2014-12-05 21:15 /user/rashmi/
nasdaq_daily_prices
Rashmis-MacBook-Pro:sqoop-1.4.5.bin__hadoop-2.0.4-alpha rashmi$ hadoop fs -ls /user/
rashmi/nasdaq_daily_prices
14/12/05 21:16:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  3 rashmi supergroup      0 2014-12-05 21:15 /user/rashmi/
nasdaq_daily_prices/_SUCCESS
-rw-r--r--  3 rashmi supergroup  538464 2014-12-05 21:15 /user/rashmi/
nasdaq_daily_prices/part-m-00000
```

Run MapReduce program on the dataset

1. Run NYSE.jar MapReduce program on the input data. input is at location “/user/rashmi/nasdaq_daily_prices/part-m-00000 /” and the output is generated at the location “/user/rashmi/nasdaq_daily_prices/output”

```
Rashmis-MacBook-Pro:Desktop rashmi$ hadoop jar NYSE.jar
com.rashmi.mapreduce.NYSE.NYSEJob /user/rashmi/nasdaq_daily_prices/part-m-00000 /
user/rashmi/nasdaq_daily_prices/output
14/12/08 16:00:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
```

14/12/08 16:00:07 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
14/12/08 16:00:07 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
14/12/08 16:00:08 INFO input.FileInputFormat: Total input paths to process : 1
14/12/08 16:00:08 INFO mapreduce.JobSubmitter: number of splits:1
14/12/08 16:00:08 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1417842624965_0007
14/12/08 16:00:08 INFO impl.YarnClientImpl: Submitted application application_1417842624965_0007
14/12/08 16:00:08 INFO mapreduce.Job: The url to track the job: http://Rashmis-MacBook-Pro.local:8088/proxy/application_1417842624965_0007/
14/12/08 16:00:08 INFO mapreduce.Job: Running job: job_1417842624965_0007
14/12/08 16:00:14 INFO mapreduce.Job: Job job_1417842624965_0007 running in uber mode : false
14/12/08 16:00:14 INFO mapreduce.Job: map 0% reduce 0%
14/12/08 16:00:20 INFO mapreduce.Job: map 100% reduce 0%
14/12/08 16:00:25 INFO mapreduce.Job: map 100% reduce 100%
14/12/08 16:00:26 INFO mapreduce.Job: Job job_1417842624965_0007 completed successfully
14/12/08 16:00:26 INFO mapreduce.Job: Counters: 49

File System Counters

- FILE: Number of bytes read=257659
- FILE: Number of bytes written=727285
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=538595
- HDFS: Number of bytes written=44
- HDFS: Number of read operations=6
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=2

Job Counters

- Launched map tasks=1
- Launched reduce tasks=1
- Data-local map tasks=1
- Total time spent by all maps in occupied slots (ms)=3430
- Total time spent by all reduces in occupied slots (ms)=3251
- Total time spent by all map tasks (ms)=3430
- Total time spent by all reduce tasks (ms)=3251
- Total vcore-seconds taken by all map tasks=3430
- Total vcore-seconds taken by all reduce tasks=3251
- Total megabyte-seconds taken by all map tasks=3512320

```

    Total megabyte-seconds taken by all reduce tasks=3329024
Map-Reduce Framework
    Map input records=9999
    Map output records=9999
    Map output bytes=237655
    Map output materialized bytes=257659
    Input split bytes=131
    Combine input records=0
    Combine output records=0
    Reduce input groups=3
    Reduce shuffle bytes=257659
    Reduce input records=9999
    Reduce output records=3
    Spilled Records=19998
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=62
    CPU time spent (ms)=0
    Physical memory (bytes) snapshot=0
    Virtual memory (bytes) snapshot=0
    Total committed heap usage (bytes)=358612992
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=538464
File Output Format Counters
    Bytes Written=44
Rashmis-MacBook-Pro:Desktop rashmi$

```

```

Rashmis-MacBook-Pro:zookeeper rashmi$ hadoop fs -ls /user/rashmi/nasdaq_daily_prices/
output
14/12/08 16:00:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  3 rashmi supergroup      0 2014-12-08 16:00 /user/rashmi/
nasdaq_daily_prices/output/_SUCCESS

```



```
-rw-r--r-- 3 rashmi supergroup    44 2014-12-08 16:00 /user/rashmi/
nasdaq_daily_prices/output/part-r-00000 [REDACTED]
Rashmis-MacBook-Pro:zookeeper rashmi$ hadoop fs -text /user/rashmi/
nasdaq_daily_prices/output/part-r-00000 [REDACTED]
14/12/08 16:00:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable [REDACTED]
AAME      4.9223E7 [REDACTED]
ABXA      7.62147968E8 [REDACTED]
AIPC      6.2642592E8 [REDACTED]
Rashmis-MacBook-Pro:zookeeper rashmi$
```

Exporting MapReduce output to MySQL using Sqoop

1. Create table stock_volume in nyse MySQL database.

```
create table stock_volume ( stock_symbol VARCHAR(35) ,
total_stock_volume double) engine innodb;
```

2. Grant privileges on nyse MySQL database.

```
mysql> grant all privileges on nyse.* to '@localhost ;
```

3. Execute sqoop export with all the required parameters to export MapReduce output into stock_volume table in MySQL

```
Rashmis-MacBook-Pro:zookeeper rashmi$ /Users/rashmi/sqoop-1.4.5/bin__hadoop-2.0.4-
alpha/bin/sqoop export --connect jdbc:mysql://localhost/nyse --table stock_volume --
export-dir /user/rashmi/nasdaq_daily_prices/output --username root --input-fields-
terminated-by '\t' --input-lines-terminated-by '\n' [REDACTED]
Warning: /Users/rashmi/sqoop-1.4.5/bin__hadoop-2.0.4-alpha/bin/../../accumulo does not
exist! Accumulo imports will fail. [REDACTED]
Please set $ACCUMULO_HOME to the root of your Accumulo installation. [REDACTED]
Warning: /Users/rashmi/sqoop-1.4.5/bin__hadoop-2.0.4-alpha/bin/../../zookeeper does not
exist! Accumulo imports will fail. [REDACTED]
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation. [REDACTED]
14/12/08 16:03:51 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5 [REDACTED]
14/12/08 16:03:51 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset. [REDACTED]
14/12/08 16:03:51 INFO tool.CodeGenTool: Beginning code generation [REDACTED]
```

14/12/08 16:03:51 INFO manager.SqlManager: Executing SQL statement: SELECT t.*
FROM `nasdaq_daily_prices_total_volume` AS t LIMIT 1
14/12/08 16:03:51 INFO manager.SqlManager: Executing SQL statement: SELECT t.*
FROM `nasdaq_daily_prices_total_volume` AS t LIMIT 1
14/12/08 16:03:51 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/
local/Cellar/hadoop/2.6.0/libexec
Note: /tmp/sqoop-rashmi/compile/fa2c9bdd67e402b57c339a5a6b9435f2/
nasdaq_daily_prices_total_volume.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
14/12/08 16:03:53 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-rashmi/
compile/fa2c9bdd67e402b57c339a5a6b9435f2/nasdaq_daily_prices_total_volume.jar
14/12/08 16:03:53 INFO mapreduce.ExportJobBase: Beginning export of
nasdaq_daily_prices_total_volume
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/Cellar/hadoop/2.6.0/libexec/share/hadoop/
common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/Cellar/hbase/0.98.8/libexec/lib/slf4j-
log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/Cellar/hive/0.14.0/libexec/lib/hive-jdbc-0.14.0-
standalone.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
14/12/08 16:03:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
14/12/08 16:03:53 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar
14/12/08 16:03:54 INFO Configuration.deprecation:
mapred.reduce.tasks.speculative.execution is deprecated. Instead, use
mapreduce.reduce.speculative
14/12/08 16:03:54 INFO Configuration.deprecation:
mapred.map.tasks.speculative.execution is deprecated. Instead, use
mapreduce.map.speculative
14/12/08 16:03:54 INFO Configuration.deprecation: mapred.map.tasks is deprecated.
Instead, use mapreduce.job.maps
14/12/08 16:03:54 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
14/12/08 16:03:55 INFO input.FileInputFormat: Total input paths to process : 1
14/12/08 16:03:55 INFO input.FileInputFormat: Total input paths to process : 1
14/12/08 16:03:55 INFO mapreduce.JobSubmitter: number of splits:4
14/12/08 16:03:55 INFO Configuration.deprecation:
mapred.map.tasks.speculative.execution is deprecated. Instead, use
mapreduce.map.speculative
14/12/08 16:03:55 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1417842624965_0008

14/12/08 16:03:55 INFO impl.YarnClientImpl: Submitted application
application_1417842624965_0008
14/12/08 16:03:55 INFO mapreduce.Job: The url to track the job: http://Rashmis-
MacBook-Pro.local:8088/proxy/application_1417842624965_0008/
14/12/08 16:03:55 INFO mapreduce.Job: Running job: job_1417842624965_0008
14/12/08 16:04:02 INFO mapreduce.Job: Job job_1417842624965_0008 running in uber
mode : false
14/12/08 16:04:02 INFO mapreduce.Job: map 0% reduce 0%
14/12/08 16:04:12 INFO mapreduce.Job: map 25% reduce 0%
14/12/08 16:04:13 INFO mapreduce.Job: map 100% reduce 0%
14/12/08 16:04:13 INFO mapreduce.Job: Job job_1417842624965_0008 completed
successfully
14/12/08 16:04:13 INFO mapreduce.Job: Counters: 30

File System Counters

- FILE: Number of bytes read=0
- FILE: Number of bytes written=458004
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=738
- HDFS: Number of bytes written=0
- HDFS: Number of read operations=16
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=0

Job Counters

- Launched map tasks=4
- Data-local map tasks=4
- Total time spent by all maps in occupied slots (ms)=33292
- Total time spent by all reduces in occupied slots (ms)=0
- Total time spent by all map tasks (ms)=33292
- Total vcore-seconds taken by all map tasks=33292
- Total megabyte-seconds taken by all map tasks=34091008

Map-Reduce Framework

- Map input records=3
- Map output records=3
- Input split bytes=616
- Spilled Records=0
- Failed Shuffles=0
- Merged Map outputs=0
- GC time elapsed (ms)=174
- CPU time spent (ms)=0
- Physical memory (bytes) snapshot=0
- Virtual memory (bytes) snapshot=0

Total committed heap usage (bytes)=541065216
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=0
14/12/08 16:04:13 INFO mapreduce.ExportJobBase: Transferred 738 bytes in 18.7436
seconds (39.3734 bytes/sec)
14/12/08 16:04:13 INFO mapreduce.ExportJobBase: Exported 3 records.

```
mysql> select * from stock_volume;  
+-----+-----+  
| stock_symbol | total_stock_volume |  
+-----+-----+  
| AIPC        | 626426000          |  
| ABXA        | 762148000          |  
| AAME        | 49223000           |  
+-----+-----+  
3 rows in set (0.00 sec)
```